

```

In [1]: # --- Environment Setup & Imports ---
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from astropy.table import Table
from astropy import units as u
from astropy import coordinates as coord
from astropy.cosmology import Planck18
from scipy.spatial import cKDTree
from sklearn.experimental import enable_halving_search_cv # Enable Halving search
from sklearn.model_selection import train_test_split, StratifiedKFold, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import (
    confusion_matrix, classification_report, accuracy_score,
    precision_score, recall_score, f1_score,
    ConfusionMatrixDisplay, RocCurveDisplay, PrecisionRecallDisplay, make_scorer
)
from sklearn.ensemble import RandomForestClassifier, HistGradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.kernel_approximation import Nystroem
from sklearn.svm import SVC, LinearSVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.pipeline import Pipeline
from sklearn.dummy import DummyClassifier
from sklearn.inspection import permutation_importance
import joblib
import warnings
import time # To time the process
from math import prod # To calculate grid size

# Plotting style configuration
plt.style.use('seaborn-v0_8-talk')
plt.rcParams.update({
    'font.size': 14,
    'axes.labelsize': 16,
    'xtick.labelsize': 12,
    'ytick.labelsize': 12,
    'legend.fontsize': 12,
    'figure.titlesize': 18
})

# Ignore warnings for cleaner output
warnings.filterwarnings('ignore', category=FutureWarning)
warnings.filterwarnings('ignore', category=UserWarning)

print("Libraries imported successfully.")
# Verify Planck18 cosmology is loaded
print(f"Using cosmology: {Planck18.name}")

```

Libraries imported successfully.

Using cosmology: Planck18

SDSS DR7 Galaxy Catalog

You will be using the SDSS DR7 survey results for the project. We are using the SDSS DR7 main galaxy sample, specifically a volume-limited version of this galaxy catalog. The catalog contains 120,606 galaxies.

```
In [2]: # --- 2. Load & Inspect SDSS DR7 Galaxy Catalog ---
print("--- Loading Galaxy Catalog ---")
try:
    # Read the ASCII file into an Astropy Table
    galaxy_catalog_ap = Table.read('SDSS_DR7_catlaog_vollim.dat', format='ascii')

    # Convert to pandas DataFrame
    galaxies_df = galaxy_catalog_ap.to_pandas()

    # Print basic info
    print("Galaxy Catalog Info:")
    print(f"Shape: {galaxies_df.shape}")
    print("\nFirst 5 rows:")
    print(galaxies_df.head())
    print("\nSummary Statistics:")
    print(galaxies_df.describe())

except FileNotFoundError:
    print("ERROR: 'SDSS_DR7_catlaog_vollim.dat' not found. Please ensure the file i
except Exception as e:
    print(f"An error occurred: {e}")
```

--- Loading Galaxy Catalog ---

Galaxy Catalog Info:

Shape: (120606, 5)

First 5 rows:

	ra	dec	redshift	Rgal	rabsmag
0	171.592148	-1.054439	0.077352	228.33	-20.697
1	174.536224	-1.051174	0.077710	229.37	-20.283
2	239.382782	-0.467646	0.084655	249.51	-20.718
3	239.679092	-0.448756	0.051608	153.14	-20.687
4	239.698471	-0.450346	0.051549	152.97	-20.197

Summary Statistics:

	ra	dec	redshift	Rgal \
count	120606.000000	120606.000000	120606.000000	120606.000000
mean	186.511186	25.595987	0.079558	234.479486
std	38.307971	17.690338	0.020233	58.922947
min	109.998665	-3.740457	0.000100	0.300000
25%	156.324512	10.579444	0.067876	200.750000
50%	186.757538	24.325087	0.082948	244.560000
75%	219.139755	38.607632	0.095995	282.270000
max	260.991974	69.879608	0.107000	313.900000

	rabsmag
count	120606.000000
mean	-20.651245
std	0.438362
min	-30.066000
25%	-20.907000
50%	-20.556000
75%	-20.300000
max	-20.090000

Right Ascension (RA) & Declination (Dec)

"RA" and "Dec" are common astronomical terms used to specify the position of celestial objects on the sky.

RA (Right Ascension) is like the longitude of the sky, measuring how far east an object is, while Dec (Declination) is like the latitude, measuring how far north or south an object is. Together, they pinpoint the precise location of celestial objects in the night sky, helping astronomers navigate and study them.

Redshifts

This galaxy catalog also has the estimated redshifts of each galaxy

Comoving Distance

Comoving distance (R_{gal}) is called Rgal in the catalog.

R_{gal} = Comoving distance in units of h^{-1} Mpc

(Note: The Hubble constant $H_0 = 100 \cdot h \text{ km s}^{-1} \text{ Mpc}^{-1}$)

R-band Absolute (AB) Magnitude

The catlaog also contains the AB magnitude of the R-band. This may be useful if you decide to use magnitude cuts.

```
In [3]: # --- 3. Load & Inspect Ground-Truth Void Catalog ---
print("\n--- Loading Void Catalog ---")

# --- Confirmed settings based on file inspection ---
void_catalog_path = 'V2_VIDE-nsa_v1_0_1_Planck2018_zobovoids.dat'
comment_char = '#'
actual_ra_col_name = 'ra'
actual_dec_col_name = 'dec'
actual_z_col_name = 'redshift'
actual_radius_col_name = 'radius'

try:
    print(f"Attempting to read ASCII file: {void_catalog_path}")
    # Use astropy.table.Table for ASCII.
    void_catalog_ap = Table.read(void_catalog_path, format='ascii',
                                comment=comment_char)

    # Convert to pandas DataFrame
    voids_df = void_catalog_ap.to_pandas()
    print(f"Successfully read {len(voids_df)} rows.")

    print("Selecting and renaming required columns...")
    # x y z redshift ra dec radius x1 y1 z1 x2 y2 z2 x3 y3 z3 area edge
    voids_df = voids_df[[actual_ra_col_name, actual_dec_col_name, actual_z_col_name,
                        actual_radius_col_name]]

    # Rename for consistency in the rest of the script
    voids_df.rename(columns={
        actual_ra_col_name: 'void_ra',
        actual_dec_col_name: 'void_dec',
        actual_z_col_name: 'void_z',
        actual_radius_col_name: 'void_radius_raw'
    }, inplace=True)

    # --- Unit Conversion Check ---
    print(f"Processing radius column: '{actual_radius_col_name}'")
    print("Input radius units are Mpc/h. Converting to Mpc.")
    h_factor = Planck18.h
    voids_df['void_radius_mpc'] = voids_df['void_radius_raw'] / h_factor

    # Print basic info
    print("\nVoid Catalog Info (Processed):")
    print(f"Number of voids: {len(voids_df)}")
    print("\nBasic Statistics:")
    # Display stats for the columns used going forward
    print(voids_df[['void_ra', 'void_dec', 'void_z', 'void_radius_mpc']].describe())
```

```
print("\nFirst 5 rows:")
print(voids_df.head())
```

```
--- Loading Void Catalog ---
Attempting to read ASCII file: V2_VIDE-nsa_v1_0_1_Planck2018_zobovoids.dat
Successfully read 531 rows.
Selecting and renaming required columns...
Processing radius column: 'radius'
ASSUMPTION: Assuming input radius units are Mpc/h. Converting to Mpc.
```

```
Void Catalog Info (Processed):
Number of voids: 531
```

Basic Statistics:

	void_ra	void_dec	void_z	void_radius_mpc
count	531.000000	531.000000	531.000000	531.000000
mean	182.360005	26.547775	0.083270	26.892926
std	39.416886	17.403740	0.022058	10.070820
min	113.824777	-1.228658	0.014676	14.787883
25%	146.810858	11.310757	0.070083	19.588240
50%	183.813816	25.326738	0.088526	24.049128
75%	214.716981	40.453067	0.101641	32.088390
max	257.577361	66.679157	0.109941	78.502145

First 5 rows:

	void_ra	void_dec	void_z	void_radius_raw	void_radius_mpc
0	148.835073	58.576377	0.102606	23.904054	35.329670
1	151.297597	22.018112	0.103102	26.164603	38.670711
2	168.297676	62.175321	0.100919	22.857079	33.782263
3	193.538452	18.004317	0.091079	24.167234	35.718643
4	140.140987	17.371752	0.102958	23.049027	34.065957

```
In [4]: # --- 4. Convert Coordinates -> 3D Cartesian ---
print("\n--- Converting Coordinates to Cartesian ---")

def spherical_to_cartesian(ra_deg, dec_deg, dist_mpc):
    """
    Converts spherical coordinates (RA, Dec, Distance) to 3D Cartesian coordinates.

    Args:
        ra_deg (array-like): Right Ascension in degrees.
        dec_deg (array-like): Declination in degrees.
        dist_mpc (array-like): Comoving distance in Mpc.

    Returns:
        tuple: (x, y, z) coordinates in Mpc.
    """
    ra_rad = np.deg2rad(ra_deg)
    dec_rad = np.deg2rad(dec_deg)

    x = dist_mpc * np.cos(dec_rad) * np.cos(ra_rad)
    y = dist_mpc * np.cos(dec_rad) * np.sin(ra_rad)
    z = dist_mpc * np.sin(dec_rad)

    return x, y, z
```

```

# Apply to Galaxy Catalog (using 'Rgal' which is assumed to be comoving distance in
# --- Unit Conversion Check ---
# If Rgal is in Mpc/h, convert to Mpc.
h = Planck18.h
galaxies_df['Rgal_Mpc'] = galaxies_df['Rgal'] / h # Convert from Mpc/h to Mpc

galaxies_df['gal_x'], galaxies_df['gal_y'], galaxies_df['gal_z'] = spherical_to_cartesian(
    galaxies_df['ra'],
    galaxies_df['dec'],
    galaxies_df['Rgal_Mpc'] # Use the converted distance
)
print("Added Cartesian coordinates (gal_x, gal_y, gal_z) to galaxy DataFrame.")
print(galaxies_df[['ra', 'dec', 'Rgal_Mpc', 'gal_x', 'gal_y', 'gal_z']].head())

# Apply to Void Catalog (calculate distance from redshift)
try:
    if 'voids_df' in locals(): # Check if void loading was successful
        voids_df['void_dist_mpc'] = Planck18.comoving_distance(voids_df['void_z']).to()

        voids_df['void_x'], voids_df['void_y'], voids_df['void_z'] = spherical_to_cartesian(
            voids_df['void_ra'],
            voids_df['void_dec'],
            voids_df['void_dist_mpc']
        )
        print("\nAdded Cartesian coordinates (void_x, void_y, void_z) and comoving distance to void DataFrame.")
    else:
        print("\nSkipping void coordinate conversion as void DataFrame was not loaded")
except Exception as e:
    print(f"\nAn error occurred during void coordinate conversion: {e}")

```

--- Converting Coordinates to Cartesian ---

Added Cartesian coordinates (gal_x, gal_y, gal_z) to galaxy DataFrame.

	ra	dec	Rgal_Mpc	gal_x	gal_y	gal_z
0	171.592148	-1.054439	337.466745	-333.783224	49.335559	-6.210196
1	174.536224	-1.051174	339.003843	-337.406819	32.273299	-6.219167
2	239.382782	-0.467646	368.770322	-187.808490	-317.349118	-3.009856
3	239.679092	-0.448756	226.337570	-114.261355	-195.371175	-1.772719
4	239.698471	-0.450346	226.086314	-114.068476	-195.192844	-1.777025

Added Cartesian coordinates (void_x, void_y, void_z) and comoving distance to void DataFrame.

	void_ra	void_dec	void_z	void_dist_mpc	void_radius_mpc	\
0	148.835073	58.576377	378.500584	443.554032	35.329670	
1	151.297597	22.018112	167.071918	445.644341	38.670711	
2	168.297676	62.175321	385.981661	436.443222	33.782263	
3	193.538452	18.004317	122.042106	394.844991	35.718643	
4	140.140987	17.371752	132.874997	445.037572	34.065957	

	void_x	void_y	void_z
0	-197.877990	119.673676	378.500584
1	-362.377112	198.415427	167.071918
2	-199.483204	41.319423	385.981661
3	-365.076403	-87.906262	122.042106
4	-326.039285	272.215100	132.874997

```

In [5]: # --- 5. Label Galaxies by Void Membership ---
print("\n--- Labeling Galaxies by Void Membership ---")

if 'galaxies_df' in locals() and 'voids_df' in locals():
    try:
        # Extract galaxy and void coordinates
        galaxy_coords = galaxies_df[['gal_x', 'gal_y', 'gal_z']].values
        void_center_coords = voids_df[['void_x', 'void_y', 'void_z']].values
        void_radii = voids_df['void_radius_mpc'].values

        # Build KDTree on void centers
        print("Building KDTree on void centers...")
        void_tree = cKDTree(void_center_coords)

        # Query the tree for nearest void for each galaxy
        print("Querying tree for nearest void for each galaxy...")
        dist_to_nearest_void, nearest_void_idx = void_tree.query(galaxy_coords, k=1)

        # Get the radius of the nearest void for each galaxy
        radius_of_nearest_void = void_radii[nearest_void_idx]

        # Label galaxies: inside void if distance <= void radius
        galaxies_df['dist_to_nearest_void'] = dist_to_nearest_void
        galaxies_df['nearest_void_idx'] = nearest_void_idx
        galaxies_df['radius_of_nearest_void'] = radius_of_nearest_void
        galaxies_df['is_void'] = (dist_to_nearest_void <= radius_of_nearest_void)

        # Print results
        print("\nVoid Membership Labeling Complete.")
        void_counts = galaxies_df['is_void'].value_counts()
        void_fraction = galaxies_df['is_void'].value_counts(normalize=True)

        print("\nGalaxy Counts by Void Membership:")
        print(void_counts)
        print("\nFraction of Galaxies by Void Membership:")
        print(void_fraction)

        print("\nSample of labeled galaxies:")
        print(galaxies_df[['gal_x', 'gal_y', 'gal_z', 'dist_to_nearest_void', 'radi

    except Exception as e:
        print(f"An error occurred during void labeling: {e}")
else:
    print("Skipping void labeling because galaxy or void DataFrame is missing.")

```

```

--- Labeling Galaxies by Void Membership ---
Building KDTree on void centers...
Querying tree for nearest void for each galaxy...

```

Void Membership Labeling Complete.

Galaxy Counts by Void Membership:

```

is_void
False    74336
True     46270
Name: count, dtype: int64

```

Fraction of Galaxies by Void Membership:

```

is_void
False    0.616354
True     0.383646
Name: proportion, dtype: float64

```

Sample of labeled galaxies:

	gal_x	gal_y	gal_z	dist_to_nearest_void \
0	-333.783224	49.335559	-6.210196	43.856471
1	-337.406819	32.273299	-6.219167	33.130764
2	-187.808490	-317.349118	-3.009856	22.435661
3	-114.261355	-195.371175	-1.772719	41.608877
4	-114.068476	-195.192844	-1.777025	41.795028

	radius_of_nearest_void	is_void
0	30.823867	False
1	30.823867	False
2	15.926156	False
3	32.650417	False
4	32.650417	False

```

In [6]: # --- 6. Exploratory Data Analysis (EDA) ---
print("\n--- Performing Exploratory Data Analysis ---")

if 'galaxies_df' in locals() and 'is_void' in galaxies_df.columns:
    # Ensure the column is boolean type before proceeding
    if galaxies_df['is_void'].dtype != 'bool':
        try:
            bool_map = {1: True, 0: False, '1': True, '0': False, True: True, False: False}
            # Apply mapping, keep existing booleans as is, convert others
            galaxies_df['is_void'] = galaxies_df['is_void'].apply(lambda x: bool_map[x])
            # Final explicit cast to bool, coercing errors to NaT/None might be safe
            galaxies_df['is_void'] = galaxies_df['is_void'].astype(bool)
            print(f"DEBUG: Conversion attempted. New dtype: {galaxies_df['is_void'].dtype}")
        except Exception as convert_err:
            print(f"DEBUG: Failed to convert 'is_void' to bool using map/astype: {convert_err}")

    try:
        print("DEBUG: Checking 'is_void' column right before plotting:")
        if 'is_void' in galaxies_df.columns:
            print(f"Data type: {galaxies_df['is_void'].dtype}")
            unique_vals = galaxies_df['is_void'].unique()
            print(f"Unique values: {unique_vals}")
            # Explicitly check they are boolean True/False

```



```

        is_boolean = all(isinstance(v, (bool, np.bool_)) for v in unique_vals)
        print(f" Are unique values boolean? {is_boolean}")
        if not is_boolean:
            raise TypeError("EDA plotting requires 'is_void' column to be str")
    else:
        raise KeyError("'is_void' column not found in DataFrame before plotting")

hue_column = 'is_void' # Use the original boolean column
plot_palette = {True: 'blue', False: 'red'} # Boolean keys for other plots
print(f"DEBUG: Using hue column '{hue_column}' with palette keys {list(plot_palette.keys())}")

# 1. 2D Wedge Plot (RA vs Comoving Distance)
print("Plotting Wedge Plot...")
plt.figure(figsize=(12, 8))
sns.scatterplot(data=galaxies_df, x='ra', y='Rgal_Mpc', hue=hue_column, s=50)
plt.title('Galaxy Distribution (Wedge Plot)')
plt.xlabel('Right Ascension (degrees)')
plt.ylabel('Comoving Distance (Mpc)')
handles, _ = plt.gca().get_legend_handles_labels()
plt.legend(handles=handles, title='Is Void Galaxy?', labels=['Non-Void (False)', 'Void (True)'])
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

# 2. Histograms of R-band Absolute Magnitude
print("Plotting Magnitude Histogram...")
plt.figure(figsize=(10, 6))
sns.histplot(data=galaxies_df, x='rabsmag', hue=hue_column, kde=True, common_kde=True)
plt.title('Distribution of R-band Absolute Magnitude')
plt.xlabel('R-band Absolute Magnitude (rabsmag)')
plt.ylabel('Density')
handles, _ = plt.gca().get_legend_handles_labels()
plt.legend(handles=handles, title='Is Void Galaxy?', labels=['Non-Void (False)', 'Void (True)'])
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

# 3. Class Balance Bar Chart (Simplified using pandas/matplotlib)
print("Plotting Class Balance (Simplified)...")
class_counts = galaxies_df['is_void'].value_counts().sort_index() # Sort False first
print("Class Counts:")
print(class_counts)
print("\nClass Proportions:")
print(galaxies_df['is_void'].value_counts(normalize=True).sort_index())

plt.figure(figsize=(7, 5))
# Define labels and colors based on sorted index (False, True)
labels = ['False (Non-Void)', 'True (Void)']
colors = [plot_palette[False], plot_palette[True]] # Use red for False, blue for True
bars = plt.bar(labels, class_counts.values, color=colors)

# Add counts on top of bars
for bar in bars:
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2.0, yval, int(yval), va='bottom')

```

```

plt.title('Class Balance: Void vs Non-Void Galaxies')
plt.xlabel('Is Void Galaxy?')
plt.ylabel('Number of Galaxies')
# plt.xticks need not be set again as labels were passed to plt.bar
plt.grid(True, axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

# 4. 3D Scatter Plot of a Subsample
print("Preparing 3D Scatter Plot...")
subsample_frac = 0.05
if 'random_state' not in locals(): random_state = 42
subsample_df = galaxies_df.sample(frac=subsample_frac, random_state=random_

# Map boolean True/False from the hue_column to colors for the 3D plot
colors_3d = subsample_df[hue_column].map({True: 'blue', False: 'red'})
if colors_3d.isnull().any():
    print("Warning: Some values could not be mapped to colors for 3D plot.")
    colors_3d = colors_3d.fillna('gray')

print(f"Plotting 3D Scatter Plot for {len(subsample_df)} ({subsample_frac*100}% of galaxies)")
fig = plt.figure(figsize=(12, 10))
ax = fig.add_subplot(111, projection='3d') # Ensure projection='3d' is correct

ax.scatter(subsample_df['gal_x'], subsample_df['gal_y'], subsample_df['gal_z'],
           c=colors_3d, s=5, alpha=0.6, marker='.') # Use mapped colors_3d
ax.set_title(f'3D Distribution of Galaxy Subsample ({subsample_frac*100:.0f}% of galaxies)')
ax.set_xlabel('X (Mpc)')
ax.set_ylabel('Y (Mpc)')
ax.set_zlabel('Z (Mpc)')
from matplotlib.lines import Line2D
legend_elements = [Line2D([0], [0], marker='o', color='w', label='Void (True)', linestyle='none'),
                   Line2D([0], [0], marker='o', color='w', label='Non-Void (False)', linestyle='none')]
ax.legend(handles=legend_elements, title="Galaxy Type")
plt.tight_layout()
plt.show()

except Exception as e:
    print(f"An error occurred during EDA: {e}")
    import traceback
    traceback.print_exc()
else:
    print("Skipping EDA because galaxy DataFrame ('galaxies_df') or 'is_void' column is missing")

```

--- Performing Exploratory Data Analysis ---

DEBUG: Checking 'is_void' column right before plotting:

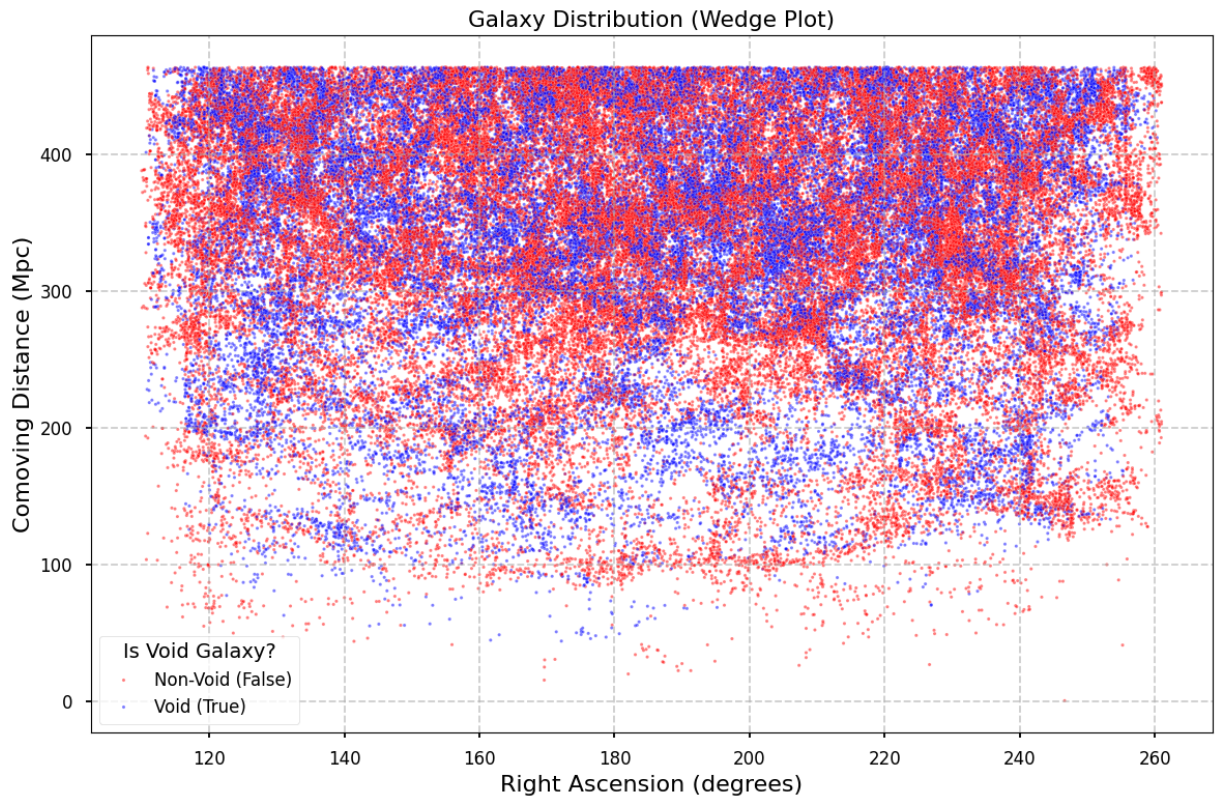
Data type: bool

Unique values: [False True]

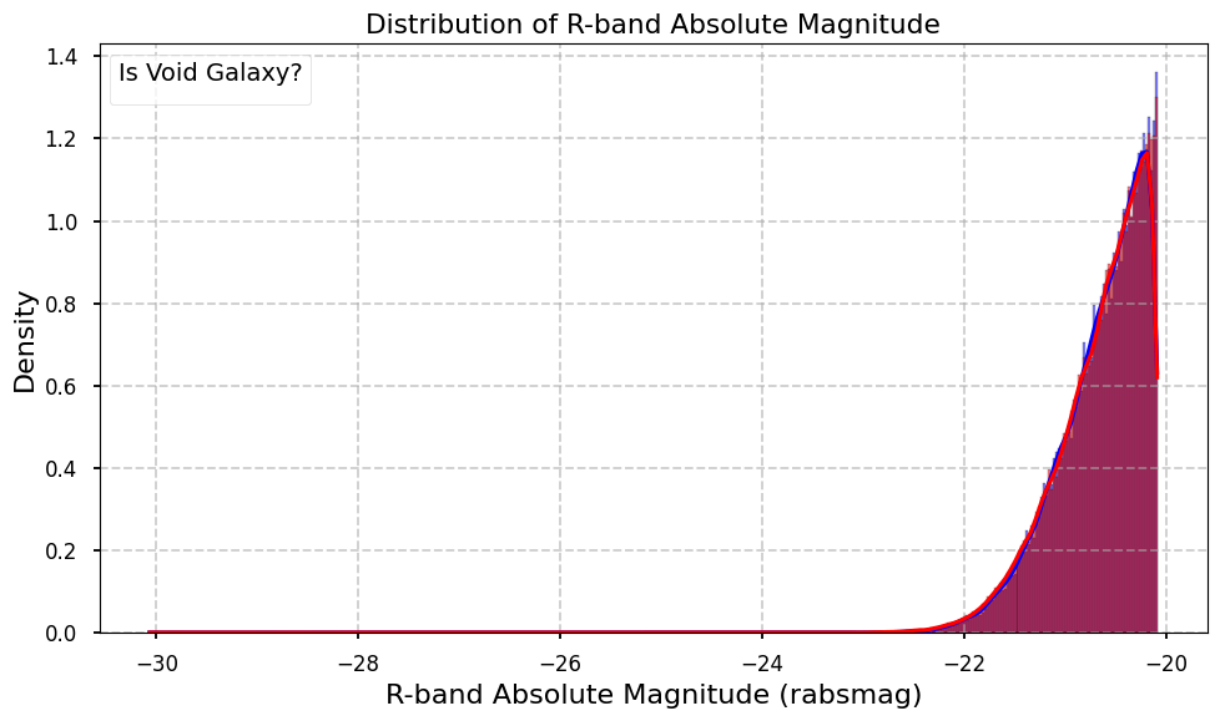
Are unique values boolean? True

DEBUG: Using hue column 'is_void' with palette keys [True, False]

Plotting Wedge Plot...



Plotting Magnitude Histogram...



Plotting Class Balance (Simplified)...

Class Counts:

is_void

False 74336

True 46270

Name: count, dtype: int64


Class Proportions:

is_void

False 0.616354

True 0.383646


Name: proportion, dtype: float64



Preparing 3D Scatter Plot...

Plotting 3D Scatter Plot for 6030 (5.0%) galaxies...

3D Distribution of Galaxy Subsample (5%)



```
In [7]: # --- 7. Feature Engineering ---
print("\n--- Performing Feature Engineering ---")

if 'galaxies_df' in locals() and 'gal_x' in galaxies_df.columns:
    try:
        # Calculate distance to the Nth nearest neighbor (NN)
        n_neighbors = 5
        print(f"Calculating distance to {n_neighbors}th nearest galaxy neighbor...")

        # Ensure we have the coordinates
        galaxy_coords = galaxies_df[['gal_x', 'gal_y', 'gal_z']].values

        # Build KDTree on galaxy coordinates
        galaxy_tree = cKDTree(galaxy_coords)

        # Query for the N+1 nearest neighbors (index 0 is the point itself)
        # We want the distance to the 5th *other* galaxy, so query for 6 neighbors.
        distances, indices = galaxy_tree.query(galaxy_coords, k=n_neighbors + 1)
```

```

# The distance to the Nth neighbor is in the Nth column (index n_neighbors)
galaxies_df['nn_dist'] = distances[:, n_neighbors]

print(f"Added '{n_neighbors}th_nn_dist' feature to galaxy DataFrame.")
print(galaxies_df[['gal_x', 'gal_y', 'gal_z', 'nn_dist']].head())

# Define feature list for scaling (will be used in the pipeline)
features_to_scale = ['gal_x', 'gal_y', 'gal_z', 'rabsmag', 'nn_dist']
print(f"\nFeatures selected for scaling: {features_to_scale}")

except Exception as e:
    print(f"An error occurred during feature engineering: {e}")
else:
    print("Skipping feature engineering because galaxy DataFrame or coordinate colu

```

```

--- Performing Feature Engineering ---
Calculating distance to 5th nearest galaxy neighbor...
Added '5th_nn_dist' feature to galaxy DataFrame.

```

	gal_x	gal_y	gal_z	nn_dist
0	-333.783224	49.335559	-6.210196	5.884006
1	-337.406819	32.273299	-6.219167	4.299565
2	-187.808490	-317.349118	-3.009856	13.771429
3	-114.261355	-195.371175	-1.772719	6.512565
4	-114.068476	-195.192844	-1.777025	6.504361

```

Features selected for scaling: ['gal_x', 'gal_y', 'gal_z', 'rabsmag', 'nn_dist']

```

```

In [8]: # --- 8. Train-Test Split & Pipeline Setup ---
print("\n--- Setting up Train-Test Split and Pipeline ---")

if 'galaxies_df' in locals() and 'is_void' in galaxies_df.columns and 'nn_dist' in
try:
    # Define features (X) and target (y)
    feature_cols = ['gal_x', 'gal_y', 'gal_z', 'rabsmag', 'nn_dist']
    X = galaxies_df[feature_cols].values
    y = galaxies_df['is_void'].astype(int).values # Target variable as integer

    print(f"Feature matrix X shape: {X.shape}")
    print(f"Target vector y shape: {y.shape}")
    print(f"Features used: {feature_cols}")
    print(f"Target variable: is_void (0=Non-Void, 1=Void)")

    # Perform stratified train-test split
    test_size = 0.20
    random_state = 42 # for reproducibility
    X_train, X_test, y_train, y_test = train_test_split(
        X, y,
        test_size=test_size,
        random_state=random_state,
        stratify=y # Important for imbalanced datasets
    )

    print(f"\nSplit data into training ({1-test_size:.0%}) and testing ({test_s
    print(f"X_train shape: {X_train.shape}, y_train shape: {y_train.shape}")
    print(f"X_test shape: {X_test.shape}, y_test shape: {y_test.shape}")

```



```

print(f"Void fraction in training set: {np.mean(y_train):.3f}")
print(f"Void fraction in test set: {np.mean(y_test):.3f}")

# Build a pipeline with scaling and a placeholder classifier
# We will replace 'classifier' during hyperparameter tuning
pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('classifier', DecisionTreeClassifier(random_state=random_state)) # Pla
])

print("\nPipeline created successfully:")
print(pipeline)

except Exception as e:
    print(f"An error occurred during train-test split or pipeline setup: {e}")
else:
    print("Skipping train-test split because required DataFrames/columns are missing")

```

--- Setting up Train-Test Split and Pipeline ---

Feature matrix X shape: (120606, 5)

Target vector y shape: (120606,)

Features used: ['gal_x', 'gal_y', 'gal_z', 'rabsmag', 'nn_dist']

Target variable: is_void (0=Non-Void, 1=Void)

Split data into training (80%) and testing (20%) sets.

X_train shape: (96484, 5), y_train shape: (96484,)

X_test shape: (24122, 5), y_test shape: (24122,)

Void fraction in training set: 0.384

Void fraction in test set: 0.384

Pipeline created successfully:

```

Pipeline(steps=[('scaler', StandardScaler()),
                 ('classifier', DecisionTreeClassifier(random_state=42))])

```

In [9]: # --- 9. Baseline Model & Metrics ---

```

print("\n--- Evaluating Baseline Model (Default Decision Tree) ---")

```

```

if 'pipeline' in locals() and 'X_train' in locals():

```

```

    try:

```

```

        # Train the baseline pipeline (Scaler + Default Decision Tree)

```

```

        pipeline.fit(X_train, y_train)

```

```

        print("Baseline pipeline trained.")

```

```

        # Evaluate on the test set

```

```

        y_pred_baseline = pipeline.predict(X_test)

```

```

        print("\nBaseline Model Performance (Decision Tree):")

```

```

        print("Confusion Matrix:")

```

```

        cm_baseline = confusion_matrix(y_test, y_pred_baseline)

```

```

        disp_baseline = ConfusionMatrixDisplay(confusion_matrix=cm_baseline, display_labels=y_test.unique())

```

```

        disp_baseline.plot(cmap=plt.cm.Blues)

```

```

        plt.title('Baseline Decision Tree Confusion Matrix')

```

```

        plt.show()

```

```

        print("\nClassification Report:")

```

```

# target_names specify labels for 0 and 1
print(classification_report(y_test, y_pred_baseline, target_names=['Non-Void', 'Void']))

# Compare with a trivial baseline (always predict the majority class, likely Non-Void)
dummy_clf = DummyClassifier(strategy="most_frequent")
dummy_clf.fit(X_train, y_train)
y_pred_dummy = dummy_clf.predict(X_test)
dummy_accuracy = accuracy_score(y_test, y_pred_dummy)

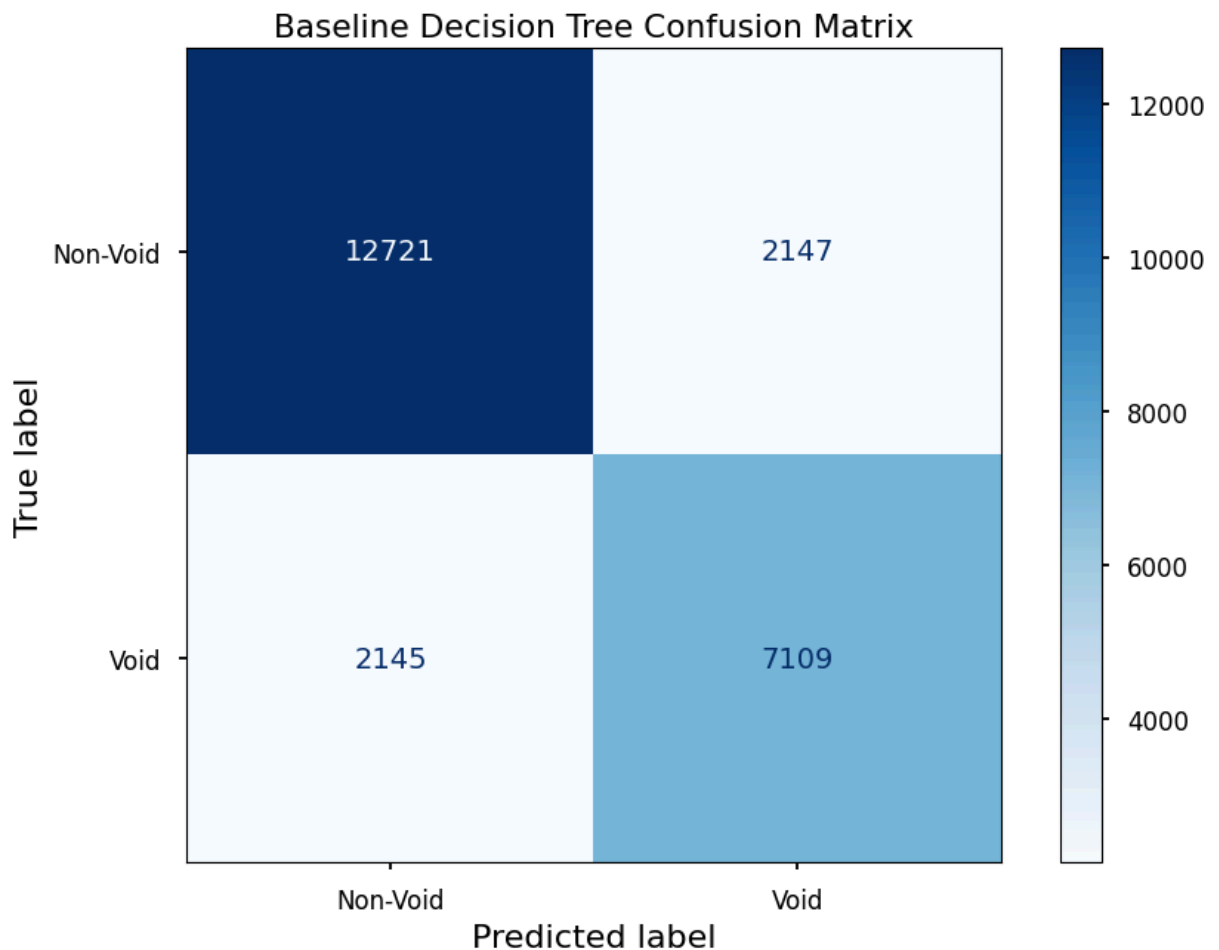
print(f"\nTrivial Baseline (Always Predict Non-Void) Accuracy: {dummy_accuracy}")
print("Note: Focus on metrics for the 'Void (1)' class (recall, precision, f1-score).")

except Exception as e:
    print(f"An error occurred during baseline model evaluation: {e}")
else:
    print("Skipping baseline evaluation because pipeline or data splits are missing")

```

--- Evaluating Baseline Model (Default Decision Tree) ---
Baseline pipeline trained.

Baseline Model Performance (Decision Tree):
Confusion Matrix:



Classification Report:

	precision	recall	f1-score	support
Non-Void (0)	0.86	0.86	0.86	14868
Void (1)	0.77	0.77	0.77	9254
accuracy			0.82	24122
macro avg	0.81	0.81	0.81	24122
weighted avg	0.82	0.82	0.82	24122

Trivial Baseline (Always Predict Non-Void) Accuracy: 0.6164

Note: Focus on metrics for the 'Void (1)' class (recall, precision, F1) as it's likely the minority class.

```
In [10]: # --- 10. Hyperparameter Tuning with Adaptive & Randomized Search ---
print("\n--- Performing Highly Optimized Hyperparameter Tuning ---")

# Check if necessary variables exist before proceeding
if 'X_train' in locals() and 'y_train' in locals() and 'random_state' in locals():
    try:
        # --- Configuration for Subsampling (ENABLED BY DEFAULT for speed) ---
        USE_SUBSAMPLE_FOR_TUNING = True # Set to True to tune on a smaller subset of
        SUBSAMPLE_SIZE = 30000 # Size of the subset if USE_SUBSAMPLE_FOR_T
        # -----

        X_tune, y_tune = X_train, y_train # Default to full training set
        if USE_SUBSAMPLE_FOR_TUNING and SUBSAMPLE_SIZE < len(X_train):
            print(f"--- NOTE: Tuning will be performed on a subsample of {SUBSAMPLE_SIZE} samples")
            # Use stratified sampling for the subset
            _, X_tune, _, y_tune = train_test_split(
                X_train, y_train,
                train_size=SUBSAMPLE_SIZE,
                stratify=y_train,
                random_state=random_state
            )
            print(f"Subsample created: X_tune shape {X_tune.shape}, y_tune shape {y_tune.shape}")
        else:
            print(f"--- NOTE: Tuning will be performed on the full training set ({len(X_train)} samples)")
            USE_SUBSAMPLE_FOR_TUNING = False

        # --- Define Classifiers including HGB and Nystroem Approx ---
        # TODO Refactor code to delete unnecessary or unused models
        print("Defining classifiers...")
        nystroem_svc_pipeline = Pipeline([
            ('nystroem', Nystroem(random_state=random_state)),
            ('linear_svc', LinearSVC(class_weight='balanced', max_iter=8000, random_state=random_state))
        ])
        # Include HGB with early stopping enabled by default in its definition
        hgb_base = HistGradientBoostingClassifier(random_state=random_state, early_stopping=True)

        classifiers = {
            'DecisionTree': DecisionTreeClassifier(random_state=random_state),
            'RandomForest': RandomForestClassifier(random_state=random_state, n_jobs=-1),
            'HistGradientBoosting': hgb_base, # Use HGB with early stopping
        }
```

```

        'KNN': KNeighborsClassifier(),
        'LinearSVC': LinearSVC(class_weight='balanced', max_iter=5000, random_s
        'RBFsvc': SVC(kernel='rbf', probability=True, class_weight='balanced',
        'NystroemApproxSVC': nystroem_svc_pipeline
    }
    tree_based_models = {"DecisionTree", "RandomForest", "HistGradientBoosting"}

    # --- Define Parameter Grids / Distributions ---
    print("Defining parameter grids/distributions...")
    param_grids = { # Use these as distributions for Randomized Search too
        'DecisionTree': {
            'classifier__max_depth': [10, 20, 30, None],
            'classifier__min_samples_leaf': [5, 10, 20, 50]
        },
        'RandomForest': {
            'classifier__n_estimators': [100, 200, 300],
            'classifier__max_depth': [10, 20, None],
            'classifier__min_samples_leaf': [1, 5, 10],
            'classifier__class_weight': ['balanced', None]
        },
        'HistGradientBoosting': {
            'classifier__max_iter': [100, 200, 300],
            'classifier__max_depth': [5, 10, 15, None],
            'classifier__learning_rate': [0.05, 0.1]
        },
        'KNN': { # Small grid - suitable for HalvingGridSearch
            'classifier__n_neighbors': [5, 11, 21],
            'classifier__weights': ['uniform', 'distance']
        },
        'LinearSVC': { # Small grid
            'classifier__C': [0.01, 0.1, 1, 10]
        },
        'RBFsvc': { # Very small grid
            'classifier__C': [1, 10],
            'classifier__gamma': [0.01, 0.1]
        },
        'NystroemApproxSVC': { # Larger grid - suitable for HalvingRandomSearch
            'classifier__nystroem__n_components': [100, 200, 300],
            'classifier__nystroem__gamma': [0.01, 0.1, 1.0],
            'classifier__linear_svc__C': [0.1, 1, 10]
        }
    }

    # --- Tuning Setup ---
    cv_folds = 3 # Use 3 folds for faster initial tuning pass
    scoring_metric = make_scorer(recall_score, pos_label=1, zero_division=0)
    print(f"Using {cv_folds}-Fold Stratified CV with HalvingGrid/RandomSearchCV")
    print(f"Optimizing for: Recall (Void Class = 1)")
    cv = StratifiedKFold(n_splits=cv_folds, shuffle=True, random_state=random_s

    best_estimators = {}
    results_summary = []
    tuning_times = {}
    RANDOM_SEARCH_THRESHOLD = 20
    N_ITER_RANDOM = 25

```

```

# --- Iterate through classifiers and tune ---
for name, classifier_obj in classifiers.items():
    start_time = time.time()
    print(f"\n--- Tuning {name} ---")

    # --- Dynamically create pipeline (skip scaler for trees) ---
    if name in tree_based_models:
        steps = [("classifier", classifier_obj)]
        print("Skipping StandardScaler for tree-based model.")
    else:
        steps = [("scaler", StandardScaler()), ("classifier", classifier_obj)]
        outer_pipeline = Pipeline(steps)
    # -----

    current_param_grid = param_grids[name]
    grid_size = prod(len(v) for v in current_param_grid.values()) if current_param_grid else 1
    print(f"Parameter grid size for {name}: {grid_size}")

    # --- Choose HalvingGridSearchCV or HalvingRandomSearchCV ---
    use_randomized = grid_size > RANDOM_SEARCH_THRESHOLD
    search_cv_args = {
        "estimator": outer_pipeline,
        "scoring": scoring_metric,
        "cv": cv,
        "factor": 3,
        "min_resources": "smallest",
        "n_jobs": -1,
        "verbose": 2,
        "random_state": random_state # Ensures reproducibility across runs/
    }

    if use_randomized:
        print(f"Using HalvingRandomSearchCV (grid size {grid_size} > {RANDOM_SEARCH_THRESHOLD})")
        search_cv = HalvingRandomSearchCV(
            **search_cv_args,
            param_distributions=current_param_grid,
            n_candidates='exhaust', # Explore specified candidates per iteration
        )
    else:
        print(f"Using HalvingGridSearchCV (grid size {grid_size} <= {RANDOM_SEARCH_THRESHOLD})")
        search_cv = HalvingGridSearchCV(
            **search_cv_args,
            param_grid=current_param_grid,
        )
    # -----

    # Fit the search object on the tuning data (full or subsample)
    print(f"Fitting {type(search_cv).__name__} for {name} using {len(y_tune)} samples")
    search_cv.fit(X_tune, y_tune)
    end_time = time.time()
    tuning_times[name] = end_time - start_time
    print(f"Finished fitting {name} in {tuning_times[name]:.2f} seconds.")

    # Store results - best estimator found by the search
    best_estimators[name] = search_cv.best_estimator_
    results_summary.append({

```

```

        'Model': name,
        'Best Score (Recall)': search_cv.best_score_,
        'Best Params': search_cv.best_params_,
        'Search Type': type(search_cv).__name__ # Record which search was used
    })

    print(f"Best {name} Recall (CV): {search_cv.best_score_:.4f}")
    print(f"Best {name} Parameters: {search_cv.best_params_}")

    if USE_SUBSAMPLE_FOR_TUNING:
        print(f"--- Refitting best {name} estimator on FULL training data")
        start_refit_time = time.time()
        try:
            # Ensure the best_estimator_ pipeline is used for refitting
            best_estimators[name].fit(X_train, y_train) # Refit on full data
            end_refit_time = time.time()
            print(f"Finished refitting in {end_refit_time - start_refit_time} seconds")
        except Exception as refit_e:
            print(f"ERROR during refitting best {name}: {refit_e}")

    # --- Display summary of tuning results ---
    print("\n--- Hyperparameter Tuning Summary ---")
    results_df = pd.DataFrame(results_summary).set_index('Model')
    results_df['Tuning Time (s)'] = results_df.index.map(tuning_times)

    # Reorder columns for clarity
    cols_order = ['Best Score (Recall)', 'Best Params', 'Search Type', 'Tuning Time (s)']
    results_df = results_df[cols_order]

    with pd.option_context('display.max_colwidth', None):
        print(results_df)

except NameError as ne:
    print(f"Skipping hyperparameter tuning because a required variable is missing: {ne}")
    import traceback; traceback.print_exc()
except ImportError:
    print("ImportError: Make sure scikit-learn is up-to-date (>= 1.0 for Halving)")
    print("Try: pip install -U scikit-learn")
    import traceback; traceback.print_exc()
except Exception as e:
    print(f"An error occurred during hyperparameter tuning: {e}")
    import traceback; traceback.print_exc()
else:
    missing_vars = []
    if 'X_train' not in locals(): missing_vars.append('X_train')
    if 'y_train' not in locals(): missing_vars.append('y_train')
    if 'random_state' not in locals(): missing_vars.append('random_state')
    print(f"Skipping hyperparameter tuning because required variables are missing: {missing_vars}")

```

```

--- Performing Highly Optimized Hyperparameter Tuning ---
--- NOTE: Tuning will be performed on a subsample of 30000 data points. ---
Subsample created: X_tune shape (66484, 5), y_tune shape (66484,)
Defining classifiers...
Defining parameter grids/distributions...
Using 3-Fold Stratified CV with HalvingGrid/RandomSearchCV.
Optimizing for: Recall (Void Class = 1)

--- Tuning DecisionTree ---
Skipping StandardScaler for tree-based model.
Parameter grid size for DecisionTree: 16
Using HalvingGridSearchCV (grid size 16 <= 20).
Fitting HalvingGridSearchCV for DecisionTree using 66484 samples...
n_iterations: 3
n_required_iterations: 3
n_possible_iterations: 8
min_resources_: 12
max_resources_: 66484
aggressive_elimination: False
factor: 3
-----
iter: 0
n_candidates: 16
n_resources: 12
Fitting 3 folds for each of 16 candidates, totalling 48 fits
-----
iter: 1
n_candidates: 6
n_resources: 36
Fitting 3 folds for each of 6 candidates, totalling 18 fits
-----
iter: 2
n_candidates: 2
n_resources: 108
Fitting 3 folds for each of 2 candidates, totalling 6 fits
Finished fitting DecisionTree in 4.04 seconds.
Best DecisionTree Recall (CV): 0.3974
Best DecisionTree Parameters: {'classifier__max_depth': None, 'classifier__min_samples_leaf': 5}
--- Refitting best DecisionTree estimator on FULL training data (96484 points)... --
-
Finished refitting in 1.04 seconds.

--- Tuning RandomForest ---
Skipping StandardScaler for tree-based model.
Parameter grid size for RandomForest: 54
Using HalvingRandomSearchCV (grid size 54 > 20).
Fitting HalvingRandomSearchCV for RandomForest using 66484 samples...
n_iterations: 4
n_required_iterations: 4
n_possible_iterations: 8
min_resources_: 12
max_resources_: 66484
aggressive_elimination: False
factor: 3
-----

```

```

iter: 0
n_candidates: 54
n_resources: 12
Fitting 3 folds for each of 54 candidates, totalling 162 fits
-----
iter: 1
n_candidates: 18
n_resources: 36
Fitting 3 folds for each of 18 candidates, totalling 54 fits
-----
iter: 2
n_candidates: 6
n_resources: 108
Fitting 3 folds for each of 6 candidates, totalling 18 fits
-----
iter: 3
n_candidates: 2
n_resources: 324
Fitting 3 folds for each of 2 candidates, totalling 6 fits
Finished fitting RandomForest in 15.39 seconds.
Best RandomForest Recall (CV): 0.4057
Best RandomForest Parameters: {'classifier__n_estimators': 100, 'classifier__min_samples_leaf': 5, 'classifier__max_depth': None, 'classifier__class_weight': 'balanced'}
--- Refitting best RandomForest estimator on FULL training data (96484 points)... --
-
Finished refitting in 4.02 seconds.

--- Tuning HistGradientBoosting ---
Skipping StandardScaler for tree-based model.
Parameter grid size for HistGradientBoosting: 24
Using HalvingRandomSearchCV (grid size 24 > 20).
Fitting HalvingRandomSearchCV for HistGradientBoosting using 66484 samples...
n_iterations: 3
n_required_iterations: 3
n_possible_iterations: 8
min_resources_: 12
max_resources_: 66484
aggressive_elimination: False
factor: 3
-----
iter: 0
n_candidates: 24
n_resources: 12
Fitting 3 folds for each of 24 candidates, totalling 72 fits
-----
iter: 1
n_candidates: 8
n_resources: 36
Fitting 3 folds for each of 8 candidates, totalling 24 fits
-----
iter: 2
n_candidates: 3
n_resources: 108
Fitting 3 folds for each of 3 candidates, totalling 9 fits
Finished fitting HistGradientBoosting in 1.46 seconds.

```

```
Best HistGradientBoosting Recall (CV): 0.4487
Best HistGradientBoosting Parameters: {'classifier__max_iter': 200, 'classifier__max_depth': None, 'classifier__learning_rate': 0.1}
--- Refitting best HistGradientBoosting estimator on FULL training data (96484 points)... ---
Finished refitting in 1.07 seconds.
```

```
--- Tuning KNN ---
Parameter grid size for KNN: 6
Using HalvingGridSearchCV (grid size 6 <= 20).
Fitting HalvingGridSearchCV for KNN using 66484 samples...
n_iterations: 2
n_required_iterations: 2
n_possible_iterations: 8
min_resources_: 12
max_resources_: 66484
aggressive_elimination: False
factor: 3
-----
iter: 0
n_candidates: 6
n_resources: 12
Fitting 3 folds for each of 6 candidates, totalling 18 fits
-----
iter: 1
n_candidates: 2
n_resources: 36
Fitting 3 folds for each of 2 candidates, totalling 6 fits
Finished fitting KNN in 0.52 seconds.
Best KNN Recall (CV): 0.2778
Best KNN Parameters: {'classifier__n_neighbors': 5, 'classifier__weights': 'distance'}
--- Refitting best KNN estimator on FULL training data (96484 points)... ---
Finished refitting in 0.10 seconds.
```

```
--- Tuning LinearSVC ---
Parameter grid size for LinearSVC: 4
Using HalvingGridSearchCV (grid size 4 <= 20).
Fitting HalvingGridSearchCV for LinearSVC using 66484 samples...
n_iterations: 2
n_required_iterations: 2
n_possible_iterations: 8
min_resources_: 12
max_resources_: 66484
aggressive_elimination: False
factor: 3
-----
iter: 0
n_candidates: 4
n_resources: 12
Fitting 3 folds for each of 4 candidates, totalling 12 fits
-----
iter: 1
n_candidates: 2
n_resources: 36
Fitting 3 folds for each of 2 candidates, totalling 6 fits
```

Finished fitting LinearSVC in 0.24 seconds.
Best LinearSVC Recall (CV): 0.5000
Best LinearSVC Parameters: {'classifier__C': 10}
--- Refitting best LinearSVC estimator on FULL training data (96484 points)... ---
Finished refitting in 0.05 seconds.

--- Tuning RBFSVC ---
Parameter grid size for RBFSVC: 4
Using HalvingGridSearchCV (grid size 4 <= 20).
Fitting HalvingGridSearchCV for RBFSVC using 66484 samples...

n_iterations: 2
n_required_iterations: 2
n_possible_iterations: 8
min_resources_: 12
max_resources_: 66484
aggressive_elimination: False
factor: 3

iter: 0
n_candidates: 4
n_resources: 12
Fitting 3 folds for each of 4 candidates, totalling 12 fits

iter: 1
n_candidates: 2
n_resources: 36
Fitting 3 folds for each of 2 candidates, totalling 6 fits
Finished fitting RBFSVC in 1087.99 seconds.

Best RBFSVC Recall (CV): 0.5556
Best RBFSVC Parameters: {'classifier__C': 10, 'classifier__gamma': 0.1}
--- Refitting best RBFSVC estimator on FULL training data (96484 points)... ---
Finished refitting in 2472.69 seconds.

--- Tuning NystroemApproxSVC ---
Parameter grid size for NystroemApproxSVC: 27
Using HalvingRandomSearchCV (grid size 27 > 20).
Fitting HalvingRandomSearchCV for NystroemApproxSVC using 66484 samples...

n_iterations: 4
n_required_iterations: 4
n_possible_iterations: 8
min_resources_: 12
max_resources_: 66484
aggressive_elimination: False
factor: 3

iter: 0
n_candidates: 27
n_resources: 12
Fitting 3 folds for each of 27 candidates, totalling 81 fits

iter: 1
n_candidates: 9
n_resources: 36
Fitting 3 folds for each of 9 candidates, totalling 27 fits

iter: 2


```

n_candidates: 3
n_resources: 108
Fitting 3 folds for each of 3 candidates, totalling 9 fits
-----
iter: 3
n_candidates: 1
n_resources: 324
Fitting 3 folds for each of 1 candidates, totalling 3 fits
Finished fitting NystroemApproxSVC in 26.87 seconds.
Best NystroemApproxSVC Recall (CV): 0.4380
Best NystroemApproxSVC Parameters: {'classifier__nystroem__n_components': 200, 'classifier__nystroem__gamma': 0.1, 'classifier__linear_svc__C': 10}
--- Refitting best NystroemApproxSVC estimator on FULL training data (96484 points)... ---
Finished refitting in 29.69 seconds.

```

```

--- Hyperparameter Tuning Summary ---
                        Best Score (Recall) \

```

Model	
DecisionTree	0.397436
RandomForest	0.405681
HistGradientBoosting	0.448718
KNN	0.277778
LinearSVC	0.500000
RBF SVC	0.555556
NystroemApproxSVC	0.438018

```

Best Params \

```

Model	
DecisionTree	{'classifier__max_depth': None, 'classifier__min_samples_leaf': 5}
RandomForest	{'classifier__n_estimators': 100, 'classifier__min_samples_leaf': 5, 'classifier__max_depth': None, 'classifier__class_weight': 'balanced'}
HistGradientBoosting	{'classifier__max_iter': 200, 'classifier__max_depth': None, 'classifier__learning_rate': 0.1}
KNN	{'classifier__n_neighbors': 5, 'classifier__weights': 'distance'}
LinearSVC	{'classifier__C': 10}
RBF SVC	{'classifier__C': 10, 'classifier__gamma': 0.1}
NystroemApproxSVC	{'classifier__nystroem__n_components': 200, 'classifier__nystroem__gamma': 0.1, 'classifier__linear_svc__C': 10}

Model	Search Type	Tuning Time (s)
DecisionTree	HalvingGridSearchCV	4.036848
RandomForest	HalvingRandomSearchCV	15.387128
HistGradientBoosting	HalvingRandomSearchCV	1.463760
KNN	HalvingGridSearchCV	0.518918
LinearSVC	HalvingGridSearchCV	0.236827
RBF SVC	HalvingGridSearchCV	1087.986660
NystroemApproxSVC	HalvingRandomSearchCV	26.874744

```

In [11]: # --- 11. Final Model Evaluation ---
print("\n--- Evaluating Tuned Models on Test Set ---")

if 'best_estimators' in locals() and 'X_test' in locals():
    evaluation_results = []

    for name, model in best_estimators.items():
        print(f"\n--- Evaluating Best {name} ---")
        try:
            # Make predictions on the test set
            y_pred = model.predict(X_test)
            # Get probabilities for ROC/PR curves if possible
            try:
                y_pred_proba = model.predict_proba(X_test)[:, 1] # Probability of c
                has_proba = True
            except AttributeError:
                print(f"Note: {name} does not support predict_proba, skipping ROC/P
                y_pred_proba = None
                has_proba = False

            # Calculate metrics
            accuracy = accuracy_score(y_test, y_pred)
            precision = precision_score(y_test, y_pred, pos_label=1)
            recall = recall_score(y_test, y_pred, pos_label=1)
            f1 = f1_score(y_test, y_pred, pos_label=1)

            evaluation_results.append({
                'Model': name,
                'Accuracy': accuracy,
                'Precision (Void)': precision,
                'Recall (Void)': recall,
                'F1 (Void)': f1
            })

            # Display Confusion Matrix
            cm = confusion_matrix(y_test, y_pred)
            disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Non
            disp.plot(cmap=plt.cm.Blues)
            plt.title(f'Best {name} Confusion Matrix (Test Set)')
            plt.show()

            # Display Classification Report
            print("\nClassification Report:")
            print(classification_report(y_test, y_pred, target_names=['Non-Void (0)

            # Plot ROC Curve (if probabilities are available)
            if has_proba:
                fig, ax = plt.subplots(1, 2, figsize=(14, 6))
                RocCurveDisplay.from_predictions(y_test, y_pred_proba, name=f'{name
                ax[0].plot([0, 1], [0, 1], 'k--', label='Chance Level')
                ax[0].set_title(f'{name} ROC Curve')
                ax[0].legend()
                ax[0].grid(True, linestyle='--', alpha=0.6)

```

```

        # Plot Precision-Recall Curve
        PrecisionRecallDisplay.from_predictions(y_test, y_pred_proba, name=
        # Calculate baseline precision (fraction of positive class)
        baseline_precision = np.sum(y_test == 1) / len(y_test)
        ax[1].axhline(baseline_precision, color='k', linestyle='--', label=
        ax[1].set_title(f'{name} Precision-Recall Curve')
        ax[1].legend()
        ax[1].grid(True, linestyle='--', alpha=0.6)

        plt.tight_layout()
        plt.show()

    except Exception as e:
        print(f"An error occurred during evaluation of {name}: {e}")

    # Summarize results in a DataFrame
    print("\n--- Final Model Performance Summary (Test Set) ---")
    evaluation_df = pd.DataFrame(evaluation_results)
    # Sort by desired metric, e.g., Recall or F1
    evaluation_df = evaluation_df.sort_values(by='Recall (Void)', ascending=False)
    print(evaluation_df.round(4)) # Round for readability


    # --- Select the final best model ---
    final_model_name = evaluation_df.iloc[0]['Model']
    final_model = best_estimators[final_model_name]
    print(f"\nSelected final model: {final_model_name}")

else:
    print("Skipping final evaluation because best estimators or test data are missi
    final_model = None # Ensure final_model is defined

--- Evaluating Tuned Models on Test Set ---

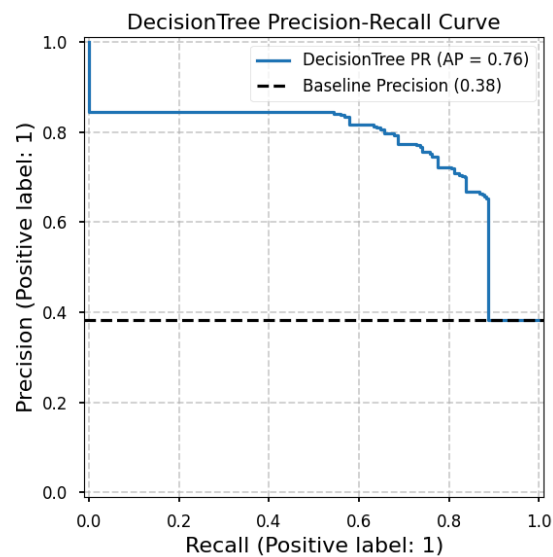
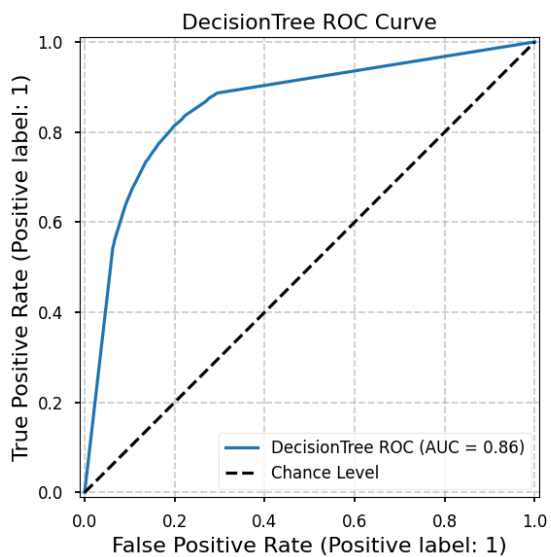
--- Evaluating Best DecisionTree ---

```




Classification Report:

	precision	recall	f1-score	support
Non-Void (0)	0.84	0.86	0.85	14868
Void (1)	0.77	0.74	0.75	9254
accuracy			0.81	24122
macro avg	0.80	0.80	0.80	24122
weighted avg	0.81	0.81	0.81	24122

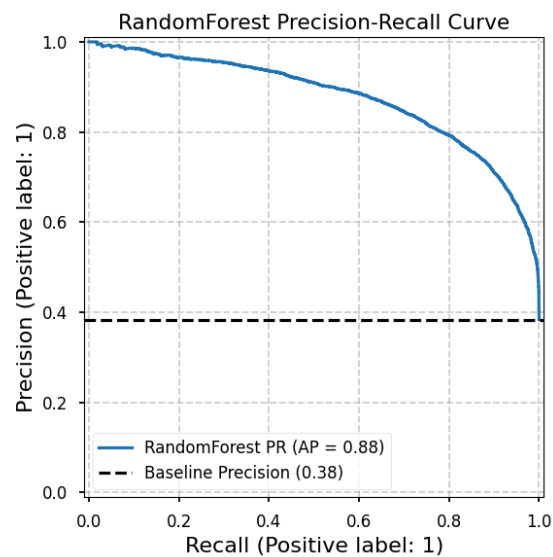



--- Evaluating Best RandomForest ---




Classification Report:

	precision	recall	f1-score	support
Non-Void (0)	0.88	0.87	0.87	14868
Void (1)	0.79	0.81	0.80	9254
accuracy			0.84	24122
macro avg	0.83	0.84	0.84	24122
weighted avg	0.84	0.84	0.84	24122

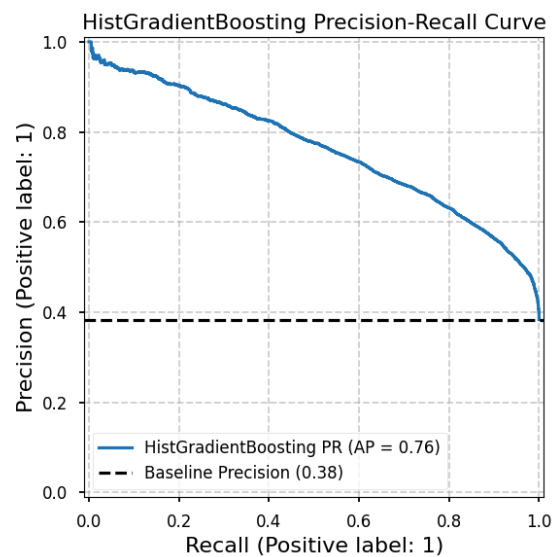
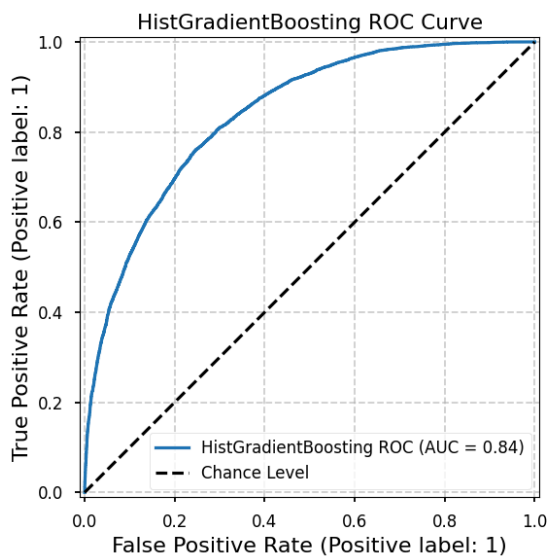


--- Evaluating Best HistGradientBoosting ---




Classification Report:

	precision	recall	f1-score	support
Non-Void (0)	0.74	0.91	0.82	14868
Void (1)	0.78	0.50	0.61	9254
accuracy			0.75	24122
macro avg	0.76	0.70	0.71	24122
weighted avg	0.76	0.75	0.74	24122

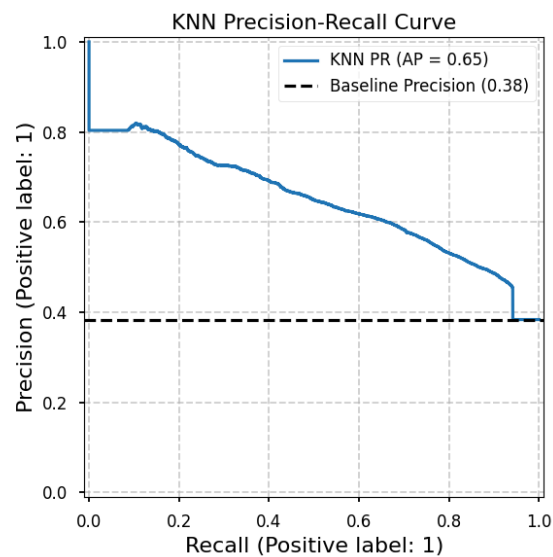
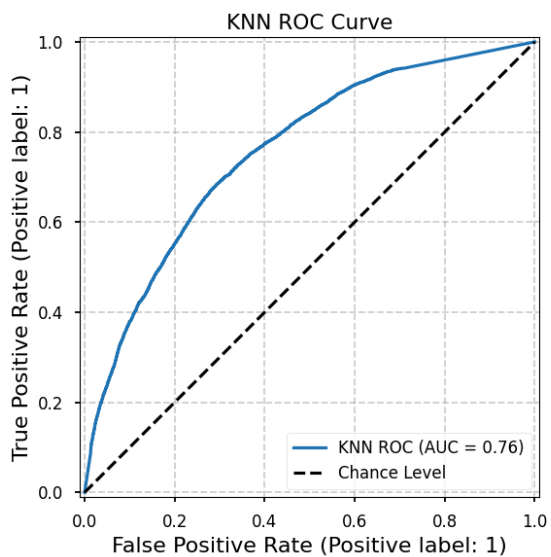


--- Evaluating Best KNN ---




Classification Report:

	precision	recall	f1-score	support
Non-Void (0)	0.74	0.80	0.77	14868
Void (1)	0.63	0.55	0.59	9254
accuracy			0.70	24122
macro avg	0.69	0.68	0.68	24122
weighted avg	0.70	0.70	0.70	24122



--- Evaluating Best LinearSVC ---


Note: LinearSVC does not support predict_proba, skipping ROC/PR curves.



Classification Report:

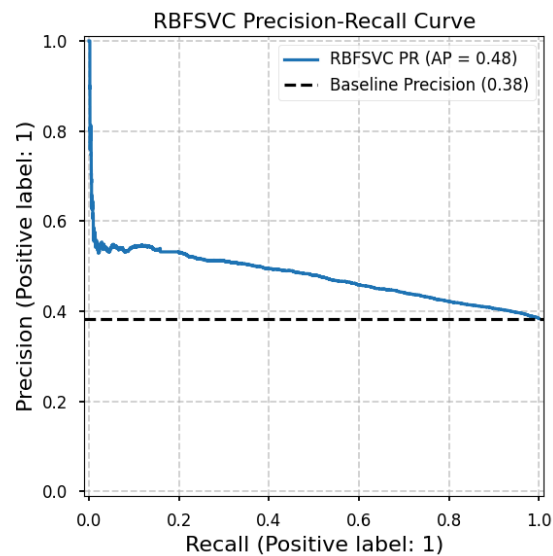
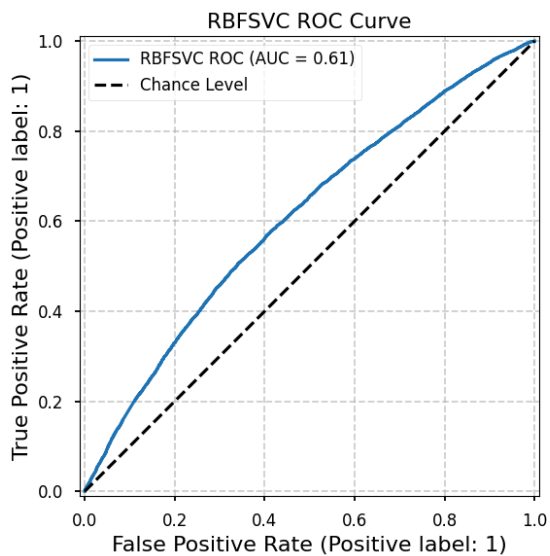
	precision	recall	f1-score	support
Non-Void (0)	0.65	0.63	0.64	14868
Void (1)	0.43	0.44	0.43	9254
accuracy			0.56	24122
macro avg	0.54	0.54	0.54	24122
weighted avg	0.56	0.56	0.56	24122

--- Evaluating Best RBFSVC ---



Classification Report:


	precision	recall	f1-score	support
Non-Void (0)	0.69	0.58	0.63	14868
Void (1)	0.46	0.58	0.52	9254
accuracy			0.58	24122
macro avg	0.58	0.58	0.57	24122
weighted avg	0.60	0.58	0.59	24122



--- Evaluating Best NystroemApproxSVC ---

Note: NystroemApproxSVC does not support predict_proba, skipping ROC/PR curves.

Best NystroemApproxSVC Confusion Matrix (Test Set)



Classification Report:

	precision	recall	f1-score	support
Non-Void (0)	0.69	0.58	0.63	14868
Void (1)	0.46	0.57	0.51	9254
accuracy			0.58	24122
macro avg	0.57	0.58	0.57	24122
weighted avg	0.60	0.58	0.58	24122

--- Final Model Performance Summary (Test Set) ---

	Model	Accuracy	Precision (Void)	Recall (Void)	F1 (Void)
1	RandomForest	0.8436	0.7910	0.8051	0.7979
0	DecisionTree	0.8136	0.7666	0.7391	0.7526
5	RBFSVC	0.5803	0.4627	0.5831	0.5160
6	NystroemApproxSVC	0.5759	0.4579	0.5739	0.5094
3	KNN	0.7050	0.6322	0.5522	0.5895
2	HistGradientBoosting	0.7521	0.7779	0.4950	0.6050
4	LinearSVC	0.5604	0.4290	0.4410	0.4349

Selected final model: RandomForest

```
In [12]: # --- 11b. Ground-Truth Comparison & Per-Void Recovery ---
print("\n--- Ground-Truth Comparison & Per-Void Recovery Analysis ---")
```

```

# Ensure we have the necessary data:
# - y_test: True labels for the test set galaxies
# - y_pred: Predictions from the chosen final model for the test set galaxies
# - galaxies_df: Original DataFrame with galaxy info, including an index that aligns
# - voids_df: DataFrame with void info (center coords, radius)
# - X_test_indices: The original indices from galaxies_df corresponding to X_test/y_test

if 'final_model' in locals() and final_model is not None and 'X_test' in locals() and 'y_test' in locals():
    try:
        # Get predictions from the final chosen model
        if 'y_pred' not in locals() or len(y_pred) != len(y_test):
            print(f"Generating predictions using the final model: {final_model_name}")
            y_pred = final_model.predict(X_test)

        # 1. Standard Classification Metrics (already printed in Step 11, but repeated here)
        print("\nOverall Test Set Performance (Final Model):")
        cm_final = confusion_matrix(y_test, y_pred)
        print("Confusion Matrix:")
        print(cm_final)
        print("\nClassification Report:")
        print(classification_report(y_test, y_pred, target_names=['Non-Void (0)', 'Void (1)'],
                                     display_labels=['Non-Void (0)', 'Void (1)'],
                                     pos_label=1) # Overall recall

        # --- Per-Void Recovery Calculation ---
        print("\nCalculating Per-Void Recovery...")

        # We need to map the test set predictions back to the original galaxy indices
        # This assumes train_test_split kept the original index if X was a DataFrame
        # or we retrieve the indices from the split. Let's assume we have the indices
        # If X was created directly from galaxies_df[feature_cols].values, the indices are lost
        # Re-run train_test_split with indices if necessary:
        _, _, _, _, idx_train, idx_test = train_test_split(
            X, y, galaxies_df.index, # Include index here
            test_size=test_size,
            random_state=random_state,
            stratify=y
        )

        # Create a temporary DataFrame for test set galaxies with their true/predicted labels
        test_galaxies_df = galaxies_df.loc[idx_test].copy()
        test_galaxies_df['true_is_void'] = y_test # Corresponds to idx_test order
        test_galaxies_df['pred_is_void'] = y_pred # Corresponds to idx_test order

        per_void_recall = []
        processed_void_indices = []
        num_test_galaxies_in_any_void = 0

        # Iterate through each *ground-truth* void
        for void_idx, void_row in voids_df.iterrows():
            # Find TRUE void galaxies (from the original labeling) that belong to this void
            # Note: A galaxy is labeled 'is_void' if it's inside its *nearest* void
            # We need galaxies in the test set whose *true* nearest void was this one
            true_members_in_test = test_galaxies_df[
                (test_galaxies_df['nearest_void_idx'] == void_idx) &
                (test_galaxies_df['true_is_void'] == 1) # Ensure they are truly void galaxies
            ]

```

```

]

if not true_members_in_test.empty:
    num_test_galaxies_in_this_void = len(true_members_in_test)
    num_test_galaxies_in_any_void += num_test_galaxies_in_this_void

    # Count how many of these were correctly predicted as void galaxies
    correctly_predicted = true_members_in_test['pred_is_void'].sum() #

    # Calculate recall for this specific void
    recall_this_void = correctly_predicted / num_test_galaxies_in_this_void
    per_void_recall.append(recall_this_void)
    processed_void_indices.append(void_idx)

# Create a DataFrame for per-void results
per_void_df = pd.DataFrame({
    'void_index': processed_void_indices,
    'recall': per_void_recall
})

# 3. Summarize per-void recall statistics
print("\nPer-Void Recall Statistics (for voids with galaxies in the test set)")
print(per_void_df['recall'].describe())

# Plot histogram of per-void recall scores
plt.figure(figsize=(10, 6))
sns.histplot(per_void_df['recall'], bins=20, kde=False)
plt.title('Histogram of Per-Void Recall Scores')
plt.xlabel('Recall (Fraction of Test Galaxies Recovered per Void)')
plt.ylabel('Number of Voids')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()

# 4. Print overall summary
mean_recall = per_void_df['recall'].mean()
std_recall = per_void_df['recall'].std()
num_voids_in_test = len(per_void_df)
total_voids = len(voids_df)

print("\n--- Per-Void Recovery Summary ---")
print(f"Analyzed {num_voids_in_test} out of {total_voids} total voids that")
print(f"Number of test-set galaxies truly belonging to these voids: {num_test_galaxies_in_test}")
print(f"Overall test set recall (Void class): {final_recall:.3f}")
print(f"Mean per-void recall: {mean_recall:.3f} (i.e., on average, recovered {mean_recall*100:.1f}% of galaxies)")
print(f"Standard deviation of per-void recall: {std_recall:.3f}")
print(f"Median per-void recall: {per_void_df['recall'].median():.3f}")

except Exception as e:
    print(f"An error occurred during per-void recovery analysis: {e}")
else:
    print("Skipping per-void recovery analysis because final model, test data, or o

```

--- Ground-Truth Comparison & Per-Void Recovery Analysis ---

Overall Test Set Performance (Final Model):

Confusion Matrix:

```
[[8580 6288]
 [3943 5311]]
```

Classification Report:


	precision	recall	f1-score	support
Non-Void (0)	0.69	0.58	0.63	14868
Void (1)	0.46	0.57	0.51	9254
accuracy			0.58	24122
macro avg	0.57	0.58	0.57	24122
weighted avg	0.60	0.58	0.58	24122

Calculating Per-Void Recovery...

Per-Void Recall Statistics (for voids with galaxies in the test set):

```
count    520.000000
mean      0.553776
std       0.358196
min       0.000000
25%       0.211722
50%       0.600000
75%       0.900000
max       1.000000
```

Name: recall, dtype: float64



--- Per-Void Recovery Summary ---

Analyzed 520 out of 531 total voids that had associated galaxies in the test set.

Number of test-set galaxies truly belonging to these voids: 9254

Overall test set recall (Void class): 0.574

Mean per-void recall: 0.554 (i.e., on average, recovered 55.4% of test galaxies per void)

Standard deviation of per-void recall: 0.358

Median per-void recall: 0.600

```
In [13]: # --- 12. Feature Importance & Interpretation ---
print("\n--- Analyzing Feature Importance ---")

# Use the selected final model
if 'final_model' in locals() and final_model is not None:
    # Check if the final model's classifier step has feature_importances_
    try:
        # Access the classifier step within the pipeline
        classifier_step = final_model.named_steps['classifier']
        feature_names = feature_cols # From Step 8

        if hasattr(classifier_step, 'feature_importances_'):
            print(f"Extracting feature importances from final model ({final_model_name})")
            importances = classifier_step.feature_importances_

            # Create DataFrame for plotting
            importance_df = pd.DataFrame({'Feature': feature_names, 'Importance': importances})
            importance_df = importance_df.sort_values(by='Importance', ascending=False)

            # Plot feature importances
            plt.figure(figsize=(10, 6))
            sns.barplot(x='Importance', y='Feature', data=importance_df, palette='vivid')
            plt.title(f'Feature Importances for {final_model_name}')
            plt.xlabel('Importance Score')
            plt.ylabel('Feature')
            plt.tight_layout()
            plt.show()

            print("\nFeature Importances:")
            print(importance_df)

        else:
            print(f"The selected model ({final_model_name}) does not have a 'feature_importances_' attribute")
            print("Consider using permutation importance for model-agnostic insight")

    # --- Permutation Importance ---
    print("\nCalculating Permutation Importance (can take time)...")
    perm_importance = permutation_importance(
        final_model, X_test, y_test, n_repeats=10, random_state=random_state,
        scoring=make_scorer(recall_score, pos_label=1) # Score based on recall
    )
    sorted_idx = perm_importance.importances_mean.argsort()[::-1] # Sort descending
    perm_importance_df = pd.DataFrame({
        'Feature': np.array(feature_names)[sorted_idx],
        'Importance Mean': perm_importance.importances_mean[sorted_idx],
        'Importance Std': perm_importance.importances_std[sorted_idx]
    })
```

```

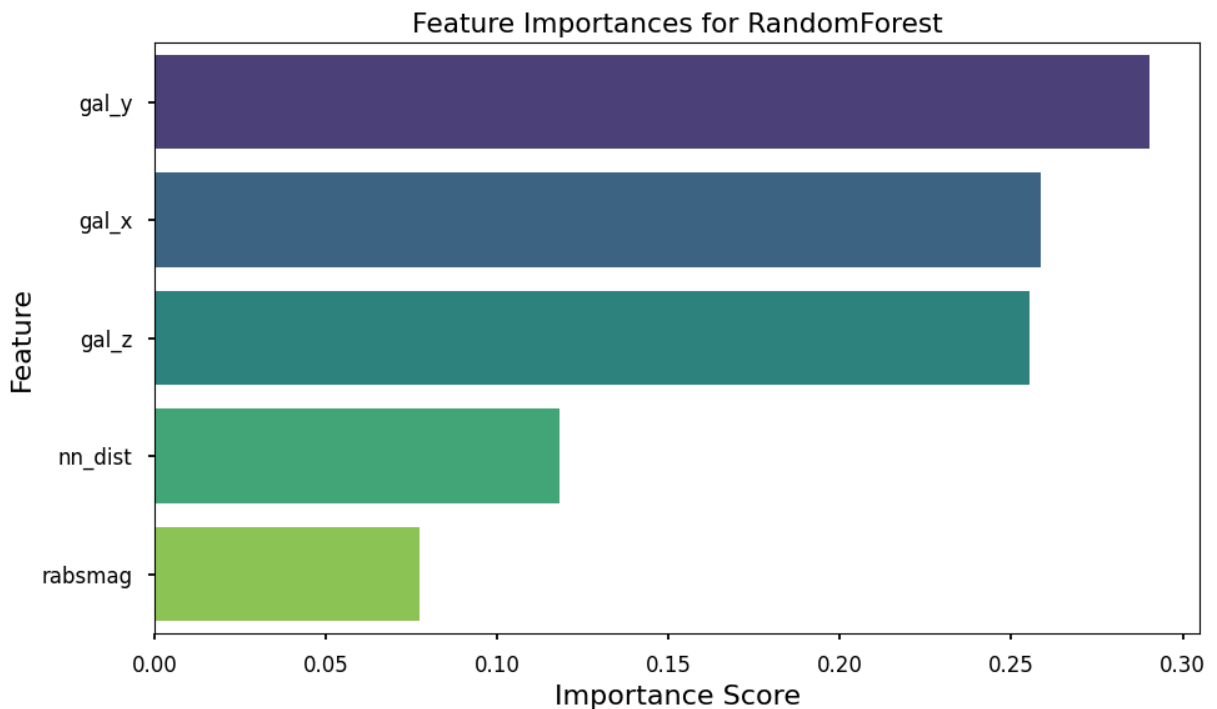
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance Mean', y='Feature', data=perm_importance_df,
plt.title(f'Permutation Importances for {final_model_name} (Scored by R
plt.xlabel('Mean Importance (Drop in Recall)')
plt.ylabel('Feature')
plt.tight_layout()
plt.show()
print("\nPermutation Importances:")
print(perm_importance_df)

except Exception as e:
    print(f"An error occurred during feature importance analysis: {e}")
else:
    print("Skipping feature importance analysis because the final model is not avai

```

--- Analyzing Feature Importance ---

Extracting feature importances from final model (RandomForest)...



Feature Importances:

	Feature	Importance
1	gal_y	0.290479
0	gal_x	0.258709
2	gal_z	0.255391
4	nn_dist	0.118181
3	rabsmag	0.077240

```

In [14]: # --- 13. 3D Visualization of Results ---
print("\n--- Creating 3D Visualization of Classification Results ---")

# Requires test set coordinates, true Labels (y_test), predicted Labels (y_pred)
# And potentially void data for overlays

if 'X_test' in locals() and 'y_test' in locals() and 'y_pred' in locals() and 'fina
try:

```

```

# Get the test set coordinates (unscaled)
# Assuming idx_test is available from Step 11b or re-run split
if 'idx_test' not in locals():
    _, _, _, _, idx_test = train_test_split(
        X, y, galaxies_df.index, test_size=test_size, random_state=random_s

test_coords_df = galaxies_df.loc[idx_test, ['gal_x', 'gal_y', 'gal_z']].copy()
test_coords_df['true_label'] = y_test
test_coords_df['pred_label'] = y_pred

# Define categories: TP, TN, FP, FN
# Void is positive (1), Non-Void is negative (0)
conditions = [
    (test_coords_df['true_label'] == 1) & (test_coords_df['pred_label'] == 1),
    (test_coords_df['true_label'] == 0) & (test_coords_df['pred_label'] == 0),
    (test_coords_df['true_label'] == 0) & (test_coords_df['pred_label'] == 1),
    (test_coords_df['true_label'] == 1) & (test_coords_df['pred_label'] == 0)
]
categories = ['True Void', 'True Non-Void', 'False Positive', 'False Negative']
colors = ['green', 'gray', 'red', 'blue'] # TP, TN, FP, FN
test_coords_df['result_category'] = np.select(conditions, categories, default='Void')
category_colors = dict(zip(categories, colors))

# --- Create 3D Plot ---
print("Generating 3D scatter plot of test set results...")
# Subsample if the test set is very large
plot_fraction = 0.2
if len(test_coords_df) * plot_fraction > 10000:
    plot_df = test_coords_df.sample(frac=plot_fraction, random_state=random_state)
    print(f"Plotting a {plot_fraction*100:.0f}% subsample ({len(plot_df)} points)")
else:
    plot_df = test_coords_df
    print(f"Plotting all {len(plot_df)} test points.")

fig = plt.figure(figsize=(14, 12))
ax = fig.add_subplot(111, projection='3d')

# Scatter plot colored by result category
scatter = ax.scatter(plot_df['gal_x'], plot_df['gal_y'], plot_df['gal_z'],
                    c=plot_df['result_category'].map(category_colors),
                    s=5, alpha=0.5, marker='o')

ax.set_title(f'3D Classification Results ({final_model_name} on Test Set)')
ax.set_xlabel('X (Mpc)')
ax.set_ylabel('Y (Mpc)')
ax.set_zlabel('Z (Mpc)')

# Create custom Legend
legend_elements = [Line2D([0], [0], marker='o', color='w', label=cat, markerfacecolor=col)
                    for cat, col in category_colors.items()]
ax.legend(handles=legend_elements, title="Result Category")
ax.grid(True)

if 'voids_df' in locals():

```



```

print("Overlaying void spheres (subset for clarity)...")
n_voids_to_plot = min(20, len(voids_df)) # Plot up to 20 voids
voids_to_plot = voids_df.nlargest(n_voids_to_plot, 'void_radius_mpc')

for _, void_row in voids_to_plot.iterrows():
    # Draw sphere wireframe
    u_sphere = np.linspace(0, 2 * np.pi, 20)
    v_sphere = np.linspace(0, np.pi, 20)
    x_sphere = void_row['void_x'] + void_row['void_radius_mpc'] * np.c
    y_sphere = void_row['void_y'] + void_row['void_radius_mpc'] * np.c
    z_sphere = void_row['void_z'] + void_row['void_radius_mpc'] * np.c
    ax.plot_wireframe(x_sphere, y_sphere, z_sphere, color='black', alp

print(f"Overlaid wireframes for {len(voids_to_plot)} largest voids.")

plt.tight_layout()
plt.show()

except Exception as e:
    print(f"An error occurred during 3D visualization: {e}")
else:
    print("Skipping 3D visualization due to missing data (test coordinates, labels,

```

--- Creating 3D Visualization of Classification Results ---


Generating 3D scatter plot of test set results...

Plotting all 24122 test points.

Overlaying void spheres (subset for clarity)...

Overlaid wireframes for 20 largest voids.

3D Classification Results (RandomForest on Test Set)



```
In [15]: # --- 14. Save Model & Predictions ---
print("\n--- Saving Final Model and Exporting Predictions ---")

if 'final_model' in locals() and final_model is not None and 'galaxies_df' in local
try:
    # 1. Save the final chosen model pipeline
    model_filename = f'final_void_classifier_{final_model_name}.joblib'
    joblib.dump(final_model, model_filename)
    print(f"Final model saved successfully to: {model_filename}")

    # 2. Add predictions for *all* galaxies to the original DataFrame
    print("Generating predictions for the *entire* galaxy dataset...")
    all_predictions = final_model.predict(galaxies_df[feature_cols].values)
    galaxies_df['pred_is_void'] = all_predictions.astype(int) # Add as integer

    print("Added 'pred_is_void' column to the main galaxy DataFrame.")
    print(galaxies_df[['is_void', 'pred_is_void']].head()) # Show true vs predi
    print("\nValue counts for predictions on full dataset:")
```

```

print(galaxies_df['pred_is_void'].value_counts(dropna=False))

# 3. Export the DataFrame to CSV
output_csv_filename = 'galaxies_with_predictions.csv'
columns_to_export = [
    'ra', 'dec', 'redshift', 'Rgal', 'Rgal_Mpc', 'rabsmag', # Original + de
    'gal_x', 'gal_y', 'gal_z', # Cartesian coords
    'nn_dist', # Engineered feature
    'dist_to_nearest_void', 'nearest_void_idx', 'radius_of_nearest_void', #
    'is_void', # True Label (ground truth)
    'pred_is_void' # Model prediction
]
# Filter out columns that might not exist if steps failed
columns_to_export = [col for col in columns_to_export if col in galaxies_df

galaxies_df[columns_to_export].to_csv(output_csv_filename, index=False)
print(f"Galaxy DataFrame with predictions exported successfully to: {output
print(f"Exported columns: {columns_to_export}")

except Exception as e:
    print(f"An error occurred during saving/exporting: {e}")
else:
    print("Skipping model saving and prediction export because the final model or g

```

--- Saving Final Model and Exporting Predictions ---

Final model saved successfully to: final_void_classifier_RandomForest.joblib

Generating predictions for the *entire* galaxy dataset...

Added 'pred_is_void' column to the main galaxy DataFrame.

	is_void	pred_is_void
0	False	0
1	False	0
2	False	0
3	False	0
4	False	0

Value counts for predictions on full dataset:

```

pred_is_void
0    73069
1    47537

```

Name: count, dtype: int64

Galaxy DataFrame with predictions exported successfully to: galaxies_with_predictions.csv

Exported columns: ['ra', 'dec', 'redshift', 'Rgal', 'Rgal_Mpc', 'rabsmag', 'gal_x', 'gal_y', 'gal_z', 'nn_dist', 'dist_to_nearest_void', 'nearest_void_idx', 'radius_of_nearest_void', 'is_void', 'pred_is_void']